

# Défis en Intelligence Artificielle

## Défi 3 : L'IA pour l'analyse et la prévision de séries temporelles (II/III)

Souhaib BEN TAIEB



December 10, 2020

# Kaggle competition

---

- Web Traffic Time Series Forecasting
  - <https://www.kaggle.com/t/17fbaf069307464094828f82a398496f>
  - **IMPORTANT:** use the previous link, **not** <https://www.kaggle.com/c/hands-on-ai-umons-2020-2021>
  - Max. five submissions per day
  - Notebooks available
- Google Colab or <https://www.kaggle.com/kernels>

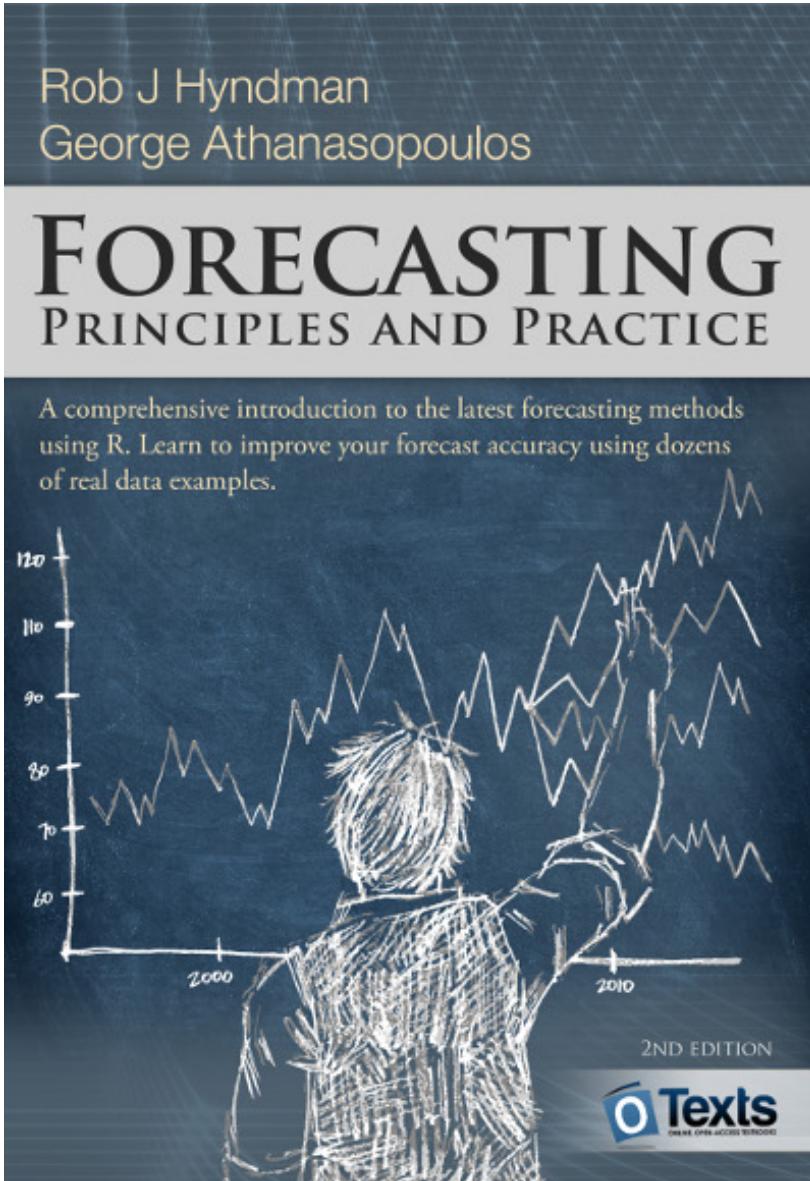
Task	Due Date	Value
Project → Kaggle submission	17 January 11:55pm	100% 35%
→ Report	24 January 11:55pm	65%

## **Part I**

# **Traditional statistical forecasting methods**

# Traditional statistical forecasting methods

---



- <https://otexts.com/fpp3/>
- Exponential smoothing methods
  - **Autoregressive integrated moving average (ARIMA)**
  - ...

# Stationarity

## Definition

If  $\{y_t\}$  is a stationary time series, then for all  $s$ , the distribution of  $(y_t, \dots, y_{t+s})$  does not depend on  $t$ .

A **stationary series** is:

- roughly horizontal
- constant variance
- no patterns predictable in the long-term

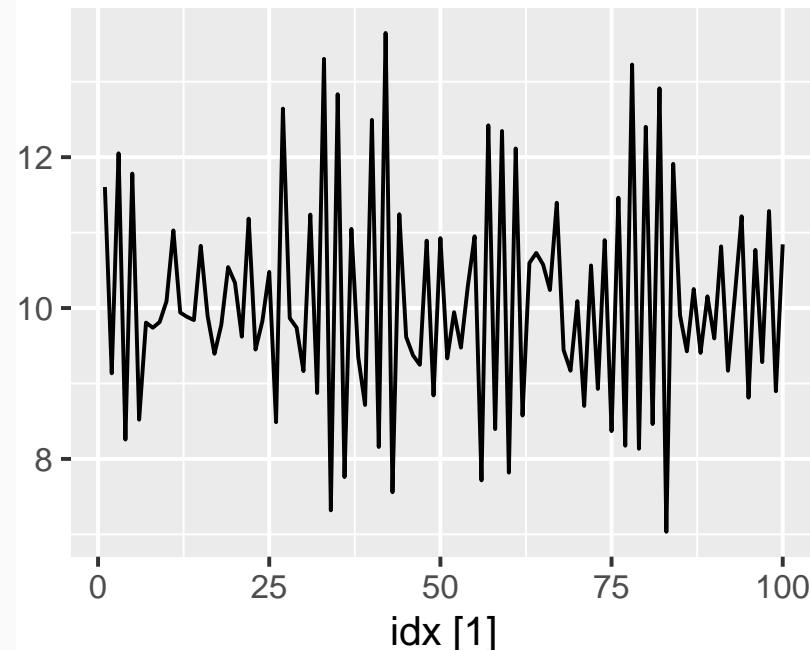
# Autoregressive models

## Autoregressive (AR) models:

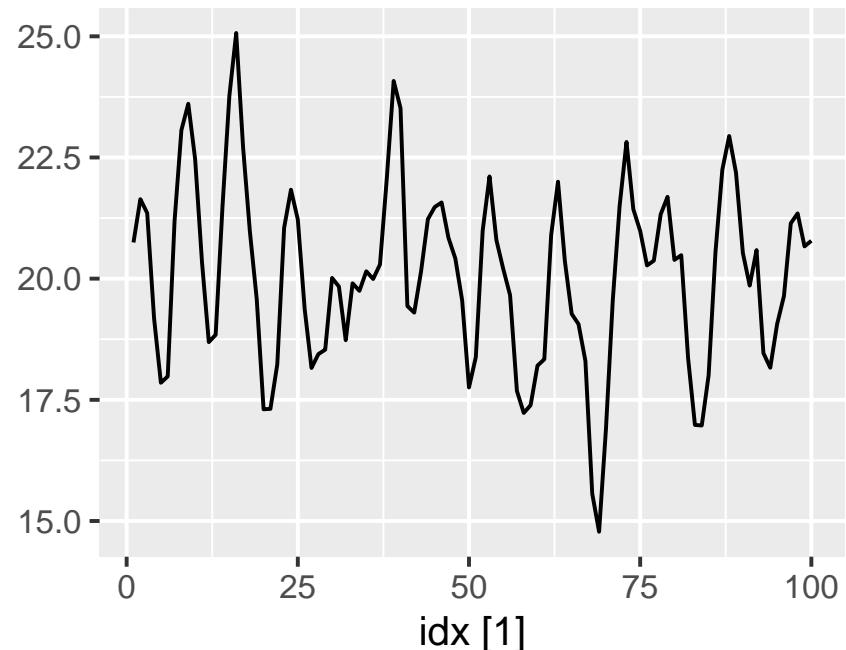
$$y_t = c + \phi_1 y_{t-1} + \phi_2 y_{t-2} + \cdots + \phi_p y_{t-p} + \varepsilon_t,$$

where  $\varepsilon_t$  is white noise. This is a multiple regression with **lagged values** of  $y_t$  as predictors.

AR(1)



AR(2)

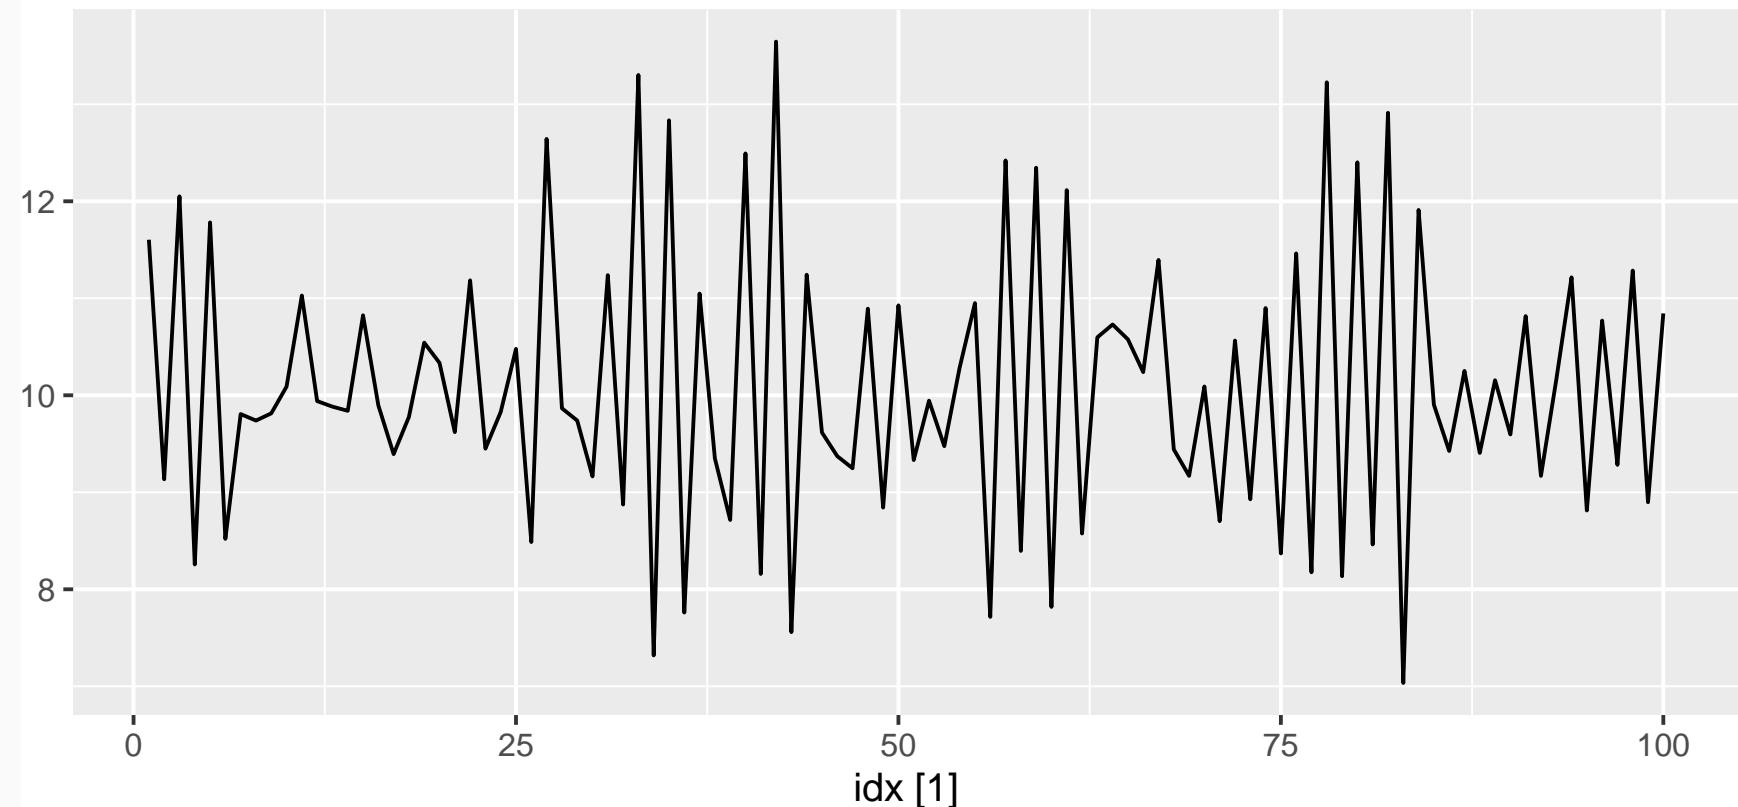


# AR(1) model

$$y_t = 18 - 0.8y_{t-1} + \varepsilon_t$$

$$\varepsilon_t \sim N(0, 1), \quad T = 100.$$

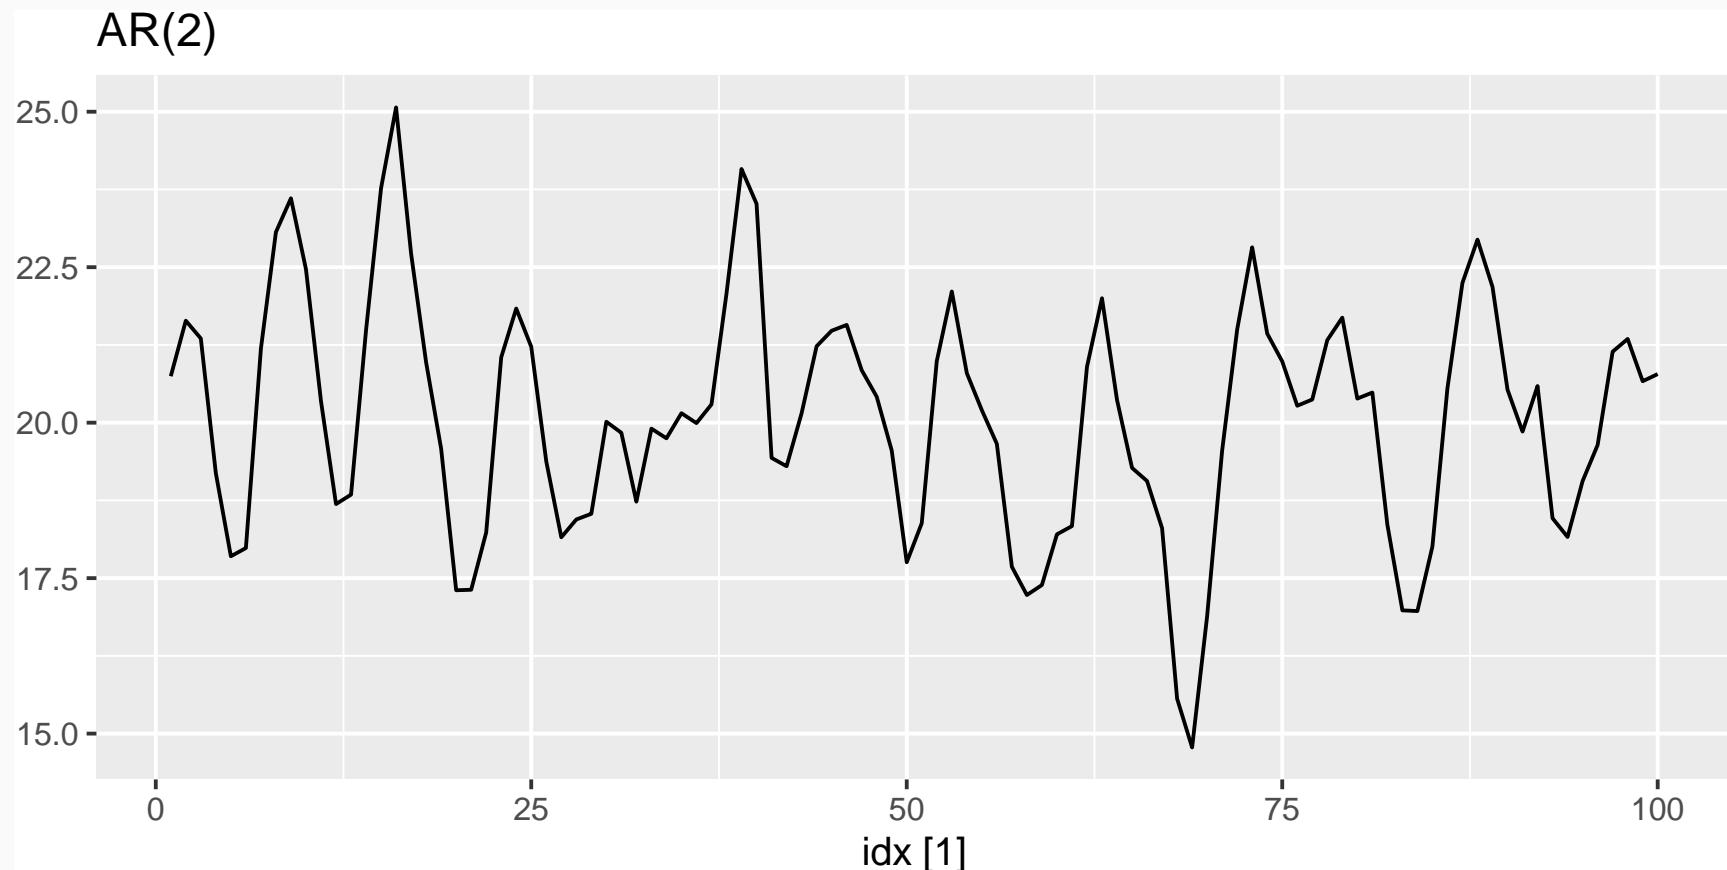
AR(1)



# AR(2) model

$$y_t = 8 + 1.3y_{t-1} - 0.7y_{t-2} + \varepsilon_t$$

$$\varepsilon_t \sim N(0, 1), \quad T = 100.$$

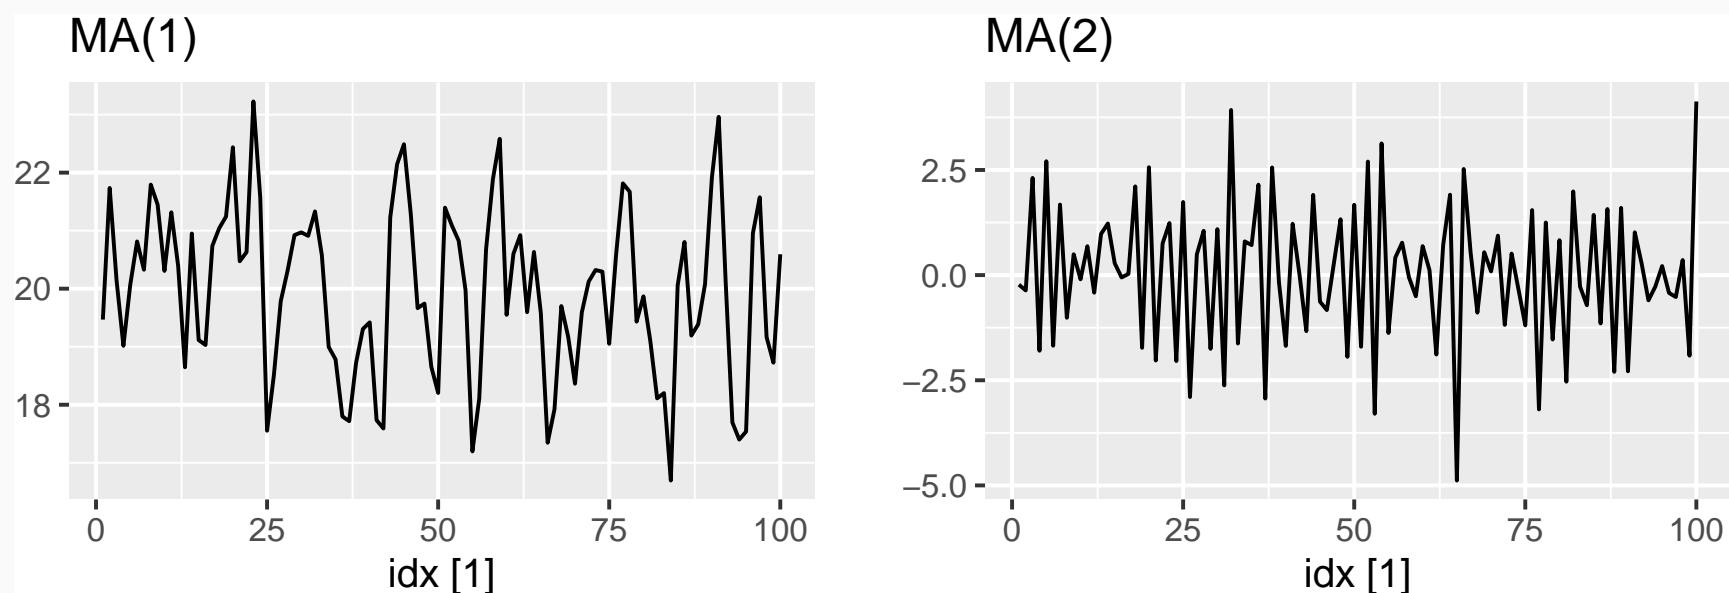


# Moving Average (MA) models

## Moving Average (MA) models:

$$y_t = c + \varepsilon_t + \theta_1 \varepsilon_{t-1} + \theta_2 \varepsilon_{t-2} + \cdots + \theta_q \varepsilon_{t-q},$$

where  $\varepsilon_t$  is white noise. This is a multiple regression with **past errors** as predictors. *Don't confuse this with moving average smoothing!*

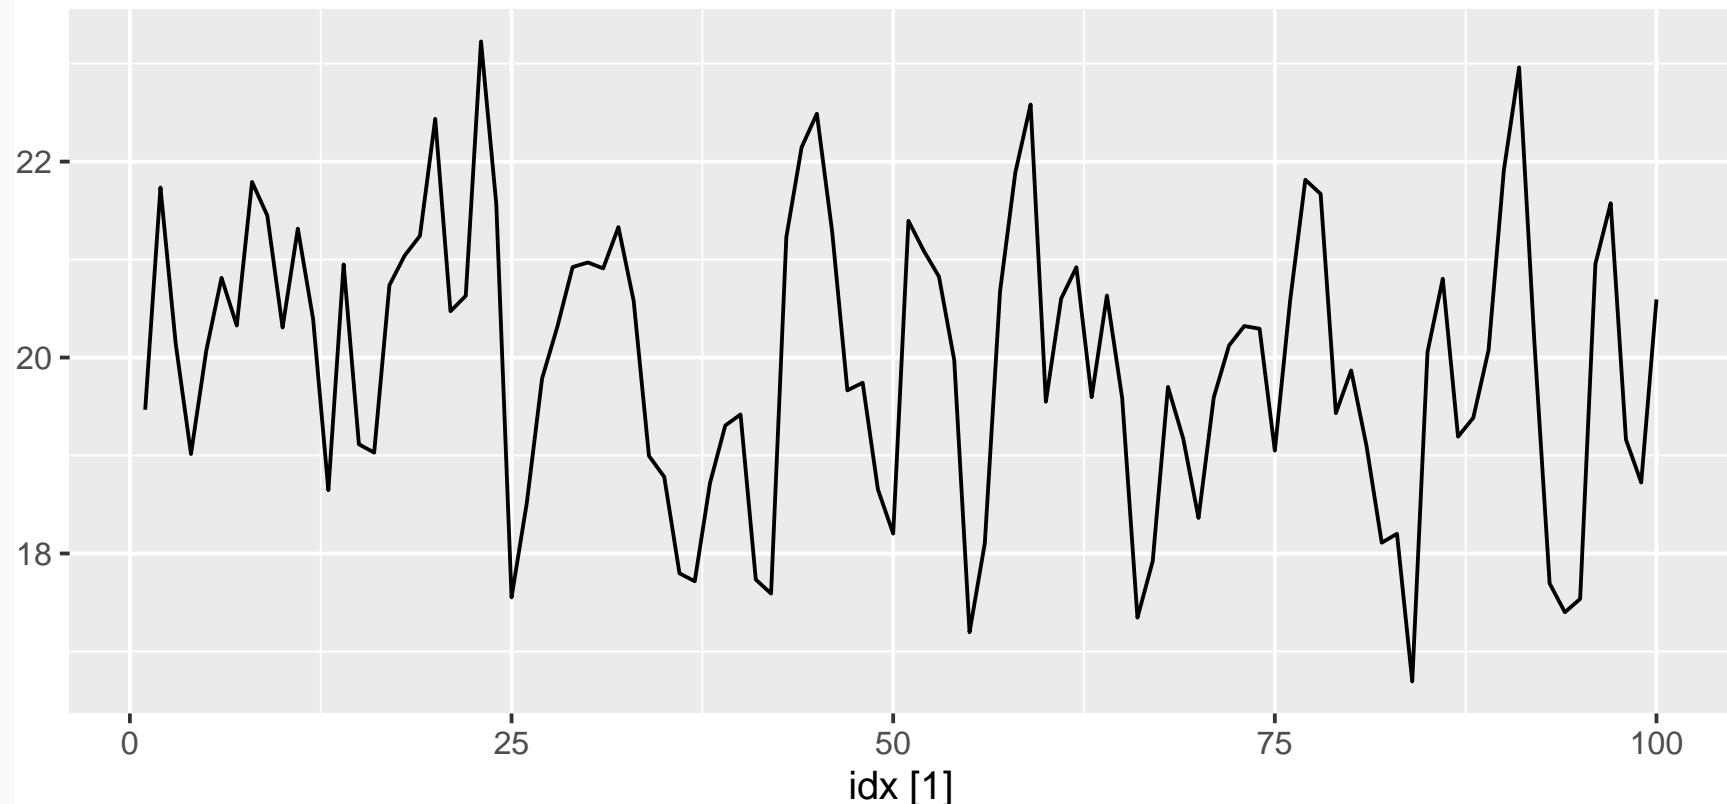


# MA(1) model

$$y_t = 20 + \varepsilon_t + 0.8\varepsilon_{t-1}$$

$$\varepsilon_t \sim N(0, 1), \quad T = 100.$$

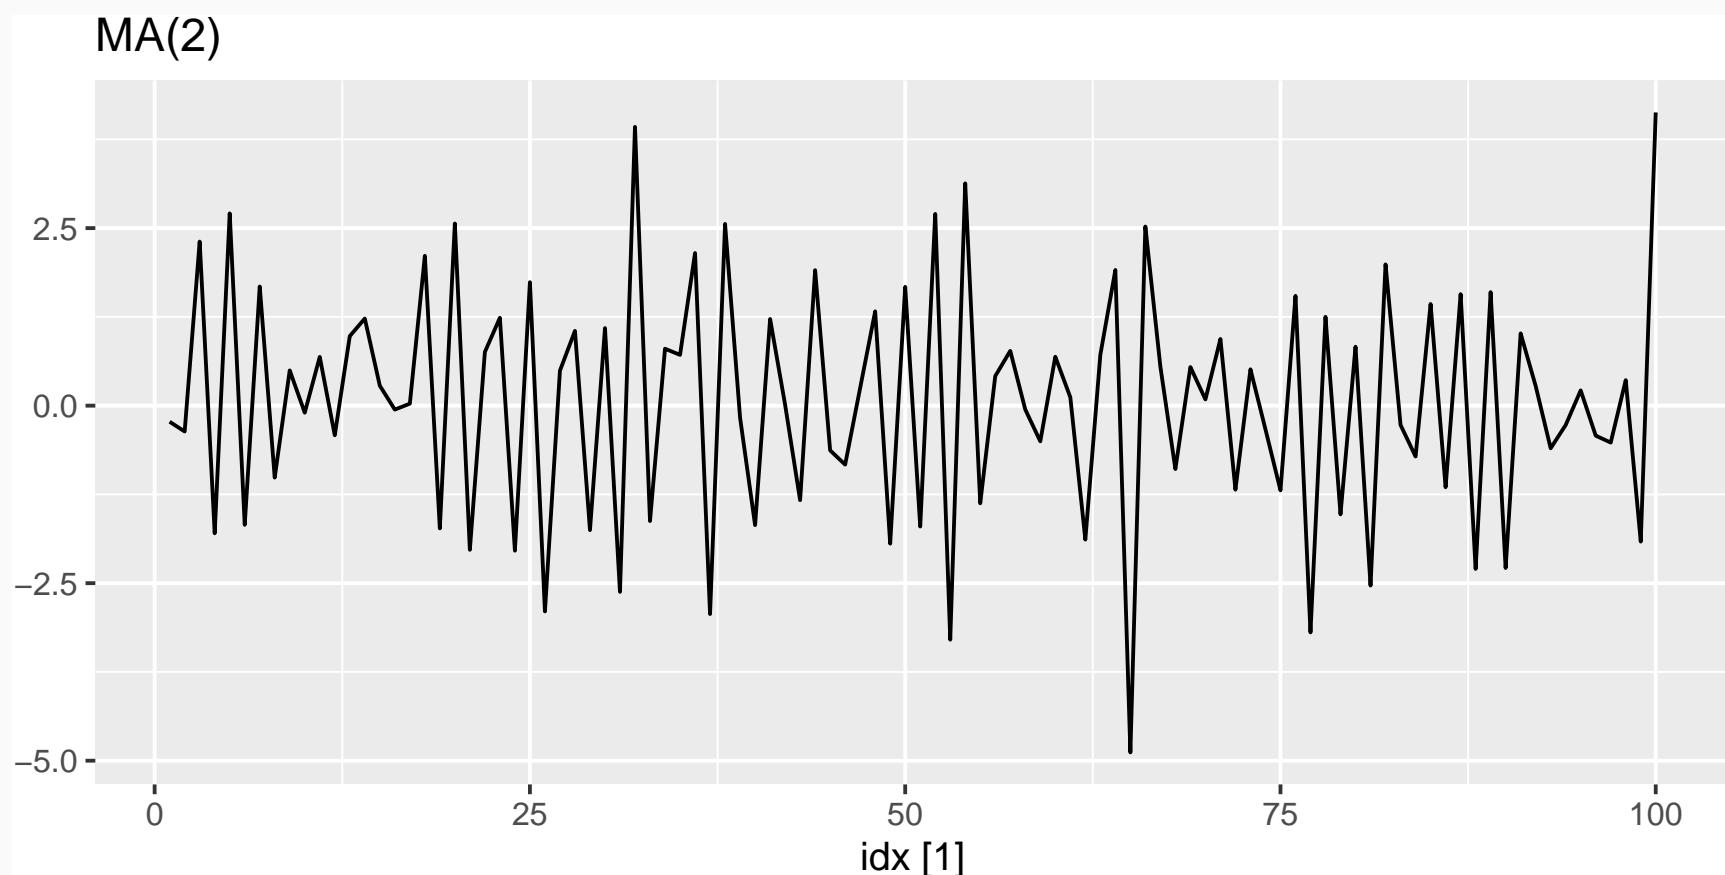
MA(1)



# MA(2) model

$$y_t = \varepsilon_t - \varepsilon_{t-1} + 0.8\varepsilon_{t-2}$$

$$\varepsilon_t \sim N(0, 1), \quad T = 100.$$



# ARMA models

## Autoregressive Moving Average models:

$$y_t = c + \phi_1 y_{t-1} + \cdots + \phi_p y_{t-p} \\ + \theta_1 \varepsilon_{t-1} + \cdots + \theta_q \varepsilon_{t-q} + \varepsilon_t.$$

- Predictors include both **lagged values of  $y_t$  and lagged errors.**
- Conditions on coefficients ensure stationarity.
-

# Maximum likelihood estimation

Having identified the model order, we need to estimate the parameters  $c, \phi_1, \dots, \phi_p, \theta_1, \dots, \theta_q$ .

- MLE is very similar to least squares estimation obtained by minimizing

$$\sum_{t=1}^T e_t^2$$

- The `auto_arima` function allows CLS or MLE estimation.
- Non-linear optimization must be used in either case.
- Different software will give different estimates.

# Information criteria

## Akaike's Information Criterion (AIC):

$$AIC = -2 \log(L) + 2(p + q + k + 1),$$

where  $L$  is the likelihood of the data,  
 $k = 1$  if  $c \neq 0$  and  $k = 0$  if  $c = 0$ .

## Corrected AIC:

$$AICc = AIC + \frac{2(p + q + k + 1)(p + q + k + 2)}{T - p - q - k - 2}.$$

## Bayesian Information Criterion:

$$BIC = AIC + [\log(T) - 2](p + q + k + 1).$$

Good models are obtained by minimizing either the AIC, AICc or BIC. Our preference is to use the AICc.

# Stationarity

## Definition

If  $\{y_t\}$  is a stationary time series, then for all  $s$ , the distribution of  $(y_t, \dots, y_{t+s})$  does not depend on  $t$ .

Transformations help to **stabilize the variance**.

For ARMA modelling, we also need to **stabilize the mean**.

# Differencing

- Differencing helps to **stabilize the mean**.
- The differenced series is the *change* between each observation in the original series:  
$$y'_t = y_t - y_{t-1}.$$
- The differenced series will have only  $T - 1$  values since it is not possible to calculate a difference  $y'_1$  for the first observation.

# Second-order differencing

Occasionally the differenced data will not appear stationary and it may be necessary to difference the data a second time:

$$\begin{aligned}y_t'' &= y'_t - y'_{t-1} \\&= (y_t - y_{t-1}) - (y_{t-1} - y_{t-2}) \\&= y_t - 2y_{t-1} + y_{t-2}.\end{aligned}$$

- $y_t''$  will have  $T - 2$  values.
- In practice, it is almost never necessary to go beyond second-order differences.

# Seasonal differencing

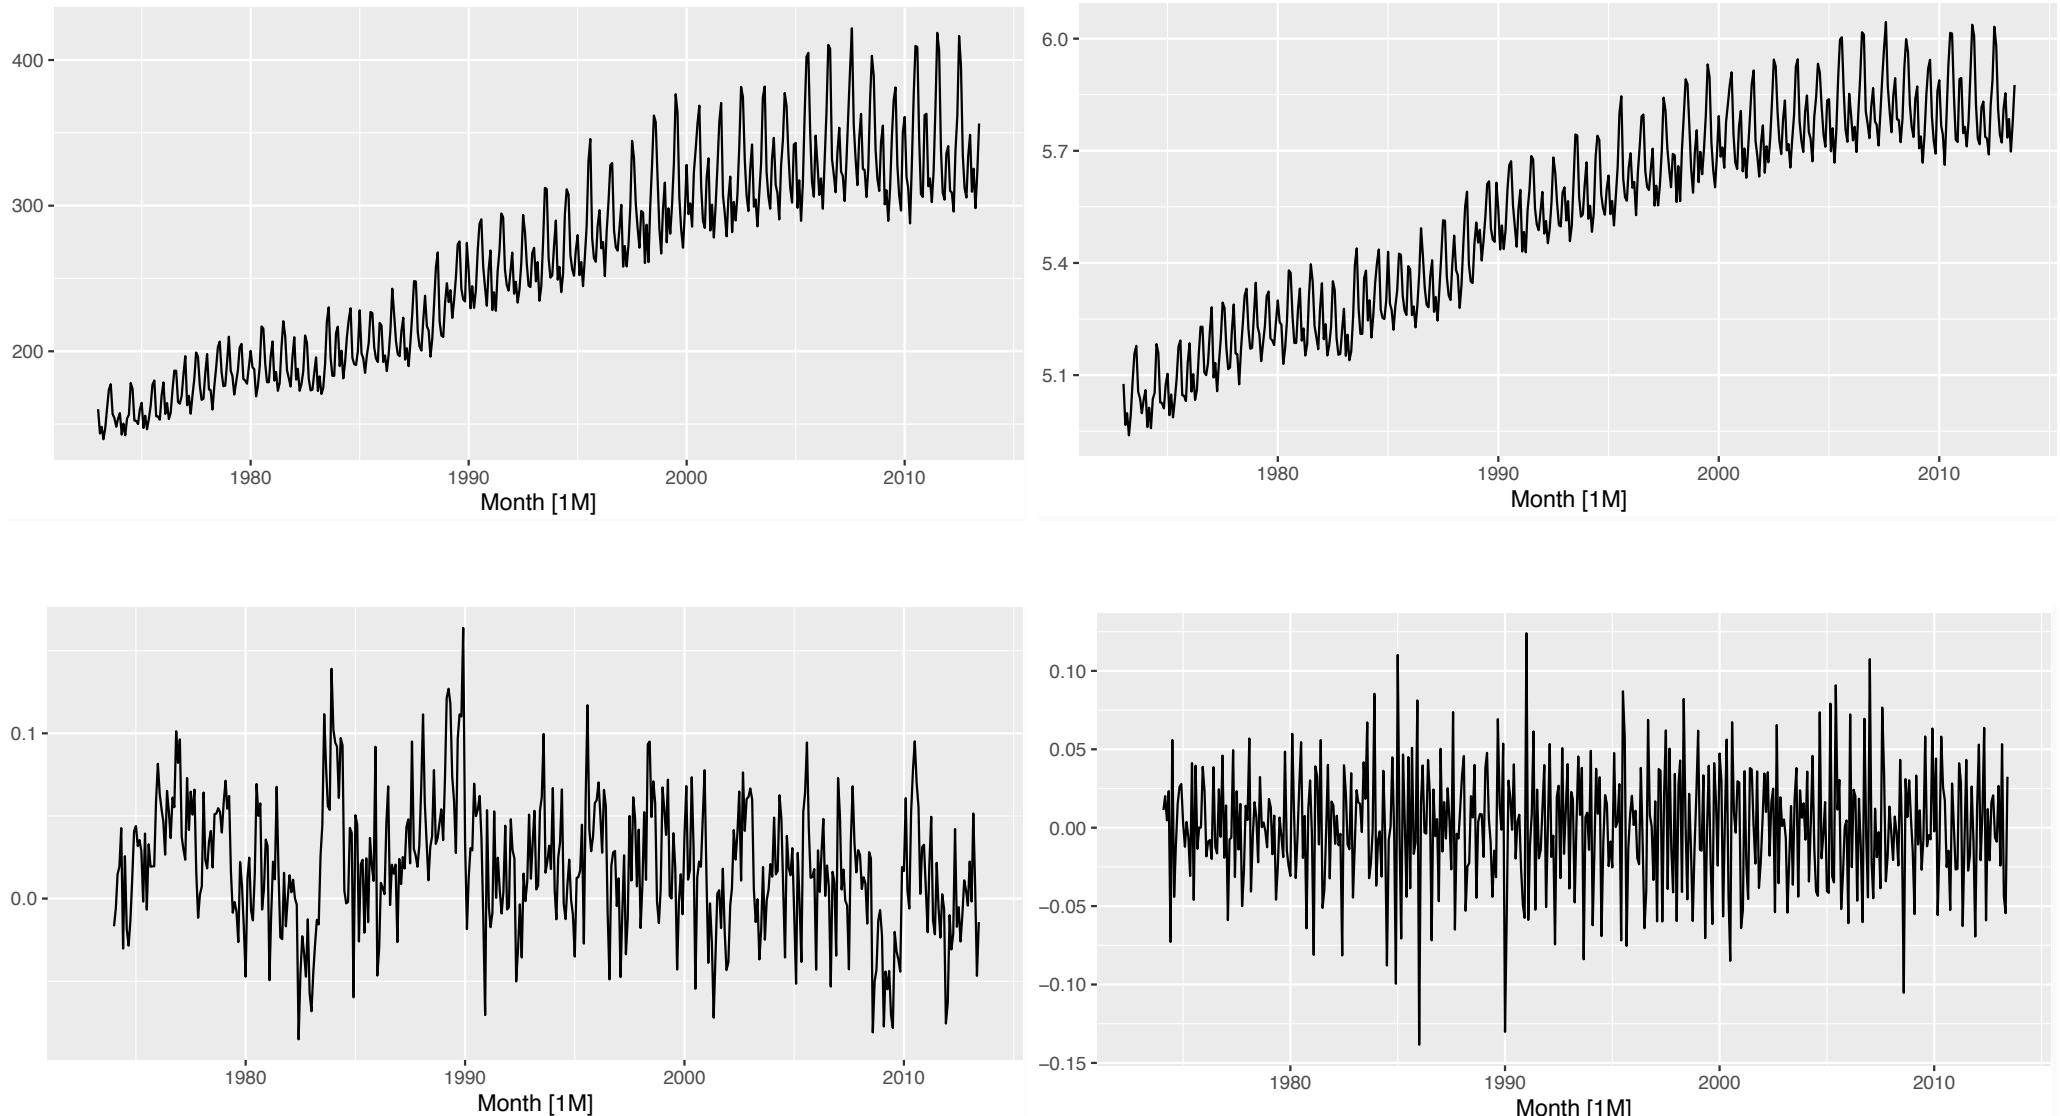
A seasonal difference is the difference between an observation and the corresponding observation from the previous year.

$$y'_t = y_t - y_{t-m}$$

where  $m$  = number of seasons.

- For monthly data  $m = 12$ .
- For quarterly data  $m = 4$ .

# Example



(1) Initial series; (2) Log transformation; (3) Seasonal difference; (4) first difference

# ARIMA models

## Autoregressive Integrated Moving Average models

### ARIMA( $p, d, q$ ) model

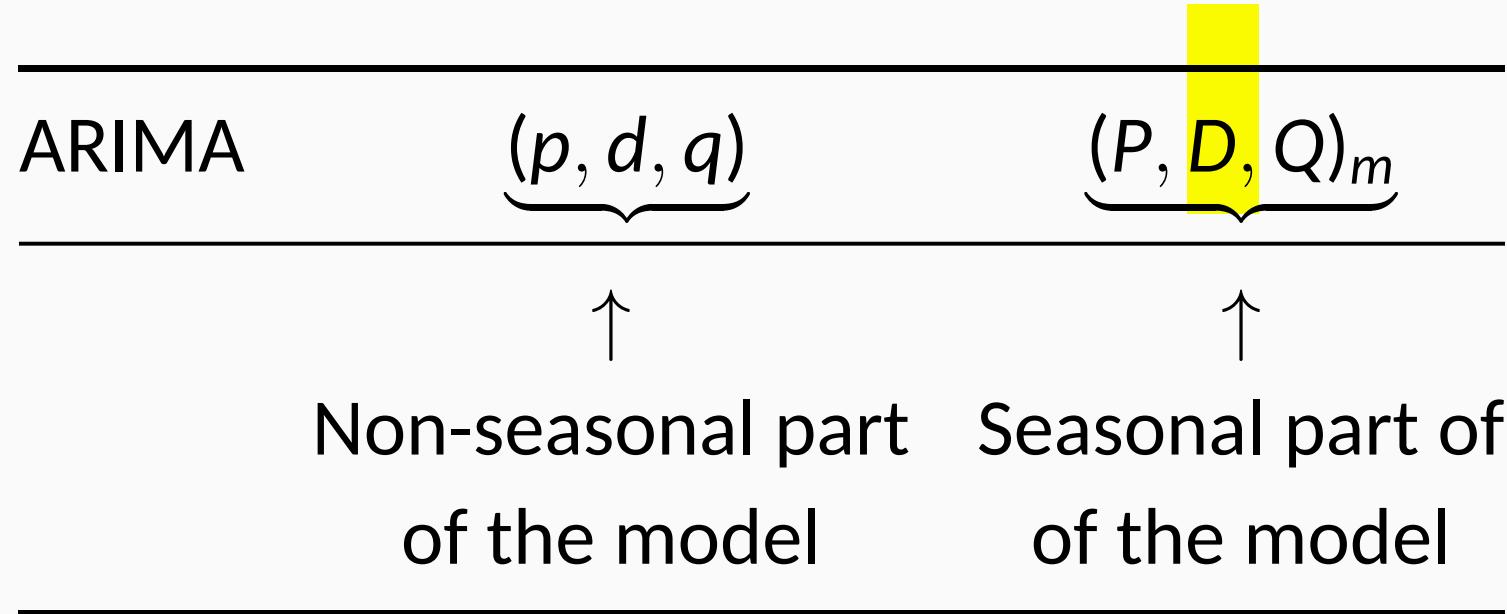
AR:  $p$  = order of the autoregressive part

I:  $d$  = degree of first differencing involved

MA:  $q$  = order of the moving average part.

- White noise model: ARIMA(0,0,0)
- Random walk: ARIMA(0,1,0) with no constant
- Random walk with drift: ARIMA(0,1,0) with const.
- AR( $p$ ): ARIMA( $p, 0, 0$ )
- MA( $q$ ): ARIMA( $0, 0, q$ )

# Seasonal ARIMA models



where  $m$  = number of observations per year.

# Software

---

## pmdarima.arima.auto\_arima

```
pmdarima.arima.auto_arima(y, X=None, start_p=2, d=None, start_q=2, max_p=5, max_d=2, max_q=5,  
start_P=1, D=None, start_Q=1, max_P=2, max_D=1, max_Q=2, max_order=5, m=1, seasonal=True,  
stationary=False, information_criterion='aic', alpha=0.05, test='kpss', seasonal_test='ocsb', stepwise=True,  
n_jobs=1, start_params=None, trend=None, method='lbfgs', maxiter=50, offset_test_args=None,  
seasonal_test_args=None, suppress_warnings=True, error_action='trace', trace=False, random=False,  
random_state=None, n_fits=10, return_valid_fits=False, out_of_sample_size=0, scoring='mse',  
scoring_args=None, with_intercept='auto', sarimax_kwds=None, **fit_args) [source] [source]
```

- <https://alkaline-ml.com/pmdarima/modules/classes.html>
- <https://alkaline-ml.com/pmdarima/modules/generated/pmdarima.arima.AutoARIMA.html#pmdarima.arima.AutoARIMA>
- [https://alkaline-ml.com/pmdarima/tips\\_and\\_tricks.html](https://alkaline-ml.com/pmdarima/tips_and_tricks.html)

# Software

SARIMAX Results						
Dep. Variable:	y	No. Observations:	1000			
Model:	SARIMAX(2, 0, 2)	Log Likelihood	-1398.466			
Date:	Thu, 10 Dec 2020	AIC	2806.931			
Time:	17:03:42	BIC	2831.470			
Sample:	0 - 1000	HQIC	2816.258			
Covariance Type:	opg					
	coef	std err	z	P> z	[0.025	0.975]
ar.L1	0.7012	0.075	9.383	0.000	0.555	0.848
ar.L2	-0.2353	0.060	-3.910	0.000	-0.353	-0.117
ma.L1	0.6982	0.072	9.724	0.000	0.557	0.839
ma.L2	0.3858	0.051	7.513	0.000	0.285	0.486
sigma2	0.9578	0.042	22.699	0.000	0.875	1.041
Ljung-Box (Q):	28.93	Jarque-Bera (JB):	3.83			
Prob(Q):	0.90	Prob(JB):	0.15			
Heteroskedasticity (H):	0.89	Skew:	0.14			
Prob(H) (two-sided):	0.28	Kurtosis:	3.09			

## Part II

# Modern statistical/machine learning forecasting methods

# AI for time series forecasting

---

- Reduce the problem of time series forecasting to one or multiple regression problems
  - Use any AI learning algorithm for regression
  - Specific AI architectures have been developed for sequential data
- Challenges
  - (Statistically) dependent data
  - Non-stationarity
  - Specific patterns: seasonality, trend, cycle, etc
  - Multi-step ahead forecasting, i.e. sequential predictions
- Model training
  - Use training data from the past to predict the future.
  - No (naive) shuffling of a time series → destroy the temporal dependence structure.

# Training with the validation set approach

---

$y_1, y_2, y_3, y_4, y_5, y_6$   $\underbrace{y_7, y_8, y_9, y_{10}}$   
Training Validation

$y_1, y_2, y_3, y_4, y_5, y_6 \longrightarrow y_7$   
 $\longrightarrow y_8$   
 $\longrightarrow y_9$   
 $\longrightarrow y_{10}$

X			y
$y_{t-2}$	$y_{t-1}$	$y_t$	$y_{t+1}$
$y_1$	$y_2$	$y_3$	$y_4$
$y_2$	$y_3$	$y_4$	$y_5$
$y_3$	$y_4$	$y_5$	$y_6$
$y_4$	$y_5$	$y_6$	$y_7$
$y_5$	$y_6$	$y_7$	$y_8$
$y_6$	$y_7$	$y_8$	$y_9$
$y_7$	$y_8$	$y_9$	$y_{10}$

# Training with the rolling-origin approach

---

$$\underbrace{y_1, y_2, y_3, y_4, y_5, y_6}_{\text{Training}}, \underbrace{y_7, y_8, y_9, y_{10}}_{\text{Validation}}$$

- $y_1, y_2, y_3, y_4, y_5, y_6 \rightarrow y_7$
- $y_1, y_2, y_3, y_4, y_5, y_6, y_7 \rightarrow y_8$
- $y_1, y_2, y_3, y_4, y_5, y_6, y_7, y_8 \rightarrow y_9$
- $y_1, y_2, y_3, y_4, y_5, y_6, y_7, y_8, y_9 \rightarrow y_{10}$

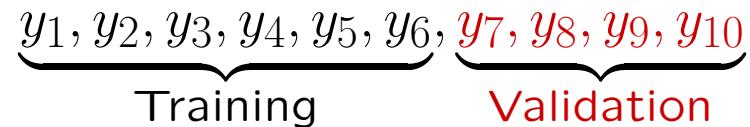
X			y
$y_{t-2}$	$y_{t-1}$	$y_t$	$y_{t+1}$
$y_1$	$y_2$	$y_3$	$y_4$
$y_2$	$y_3$	$y_4$	$y_5$
$y_3$	$y_4$	$y_5$	$y_6$
$y_4$	$y_5$	$y_6$	$y_7$

X			y
$y_{t-2}$	$y_{t-1}$	$y_t$	$y_{t+1}$
$y_1$	$y_2$	$y_3$	$y_4$
$y_2$	$y_3$	$y_4$	$y_5$
$y_3$	$y_4$	$y_5$	$y_6$
$y_4$	$y_5$	$y_6$	$y_7$
$y_5$	$y_6$	$y_7$	$y_8$

X			y
$y_{t-2}$	$y_{t-1}$	$y_t$	$y_{t+1}$
$y_1$	$y_2$	$y_3$	$y_4$
$y_2$	$y_3$	$y_4$	$y_5$
$y_3$	$y_4$	$y_5$	$y_6$
$y_4$	$y_5$	$y_6$	$y_7$
$y_5$	$y_6$	$y_7$	$y_8$
$y_6$	$y_7$	$y_8$	$y_9$

X			y
$y_{t-2}$	$y_{t-1}$	$y_t$	$y_{t+1}$
$y_1$	$y_2$	$y_3$	$y_4$
$y_2$	$y_3$	$y_4$	$y_5$
$y_3$	$y_4$	$y_5$	$y_6$
$y_4$	$y_5$	$y_6$	$y_7$
$y_5$	$y_6$	$y_7$	$y_8$
$y_6$	$y_7$	$y_8$	$y_9$
$y_7$	$y_8$	$y_9$	$y_{10}$

# Training with the rolling-origin approach



- $y_1, y_2, y_3, y_4, y_5, y_6 \longrightarrow y_7$
- $y_2, y_3, y_4, y_5, y_6, y_7 \longrightarrow y_8$
- $y_3, y_4, y_5, y_6, y_7, y_8 \longrightarrow y_9$
- $y_4, y_5, y_6, y_7, y_8, y_9 \longrightarrow y_{10}$

X			y
$y_{t-2}$	$y_{t-1}$	$y_t$	$y_{t+1}$
$y_1$	$y_2$	$y_3$	$y_4$
$y_2$	$y_3$	$y_4$	$y_5$
$y_3$	$y_4$	$y_5$	$y_6$
$y_4$	$y_5$	$y_6$	$y_7$
$y_4$	$y_5$	$y_6$	$y_7$

X			y
$y_{t-2}$	$y_{t-1}$	$y_t$	$y_{t+1}$
$y_2$	$y_3$	$y_4$	$y_5$
$y_3$	$y_4$	$y_5$	$y_6$
$y_4$	$y_5$	$y_6$	$y_7$
$y_5$	$y_6$	$y_7$	$y_8$

X			y
$y_{t-2}$	$y_{t-1}$	$y_t$	$y_{t+1}$
$y_3$	$y_4$	$y_5$	$y_6$
$y_4$	$y_5$	$y_6$	$y_7$
$y_5$	$y_6$	$y_7$	$y_8$
$y_6$	$y_7$	$y_8$	$y_9$

X			y
$y_{t-2}$	$y_{t-1}$	$y_t$	$y_{t+1}$
$y_4$	$y_5$	$y_6$	$y_7$
$y_5$	$y_6$	$y_7$	$y_8$
$y_6$	$y_7$	$y_8$	$y_9$
$y_7$	$y_8$	$y_9$	$y_{10}$

# Multi-step forecasting - recursive strategy

---

$$y_1, y_2, y_3, y_4, y_5, y_6, y_7, y_8, y_9, y_{10} \rightarrow ?, ?, ?$$

<b>X</b>			<b>y</b>
$y_{t-2}$	$y_{t-1}$	$y_t$	$y_{t+1}$
$y_1$	$y_2$	$y_3$	$y_4$
$y_2$	$y_3$	$y_4$	$y_5$
$y_3$	$y_4$	$y_5$	$y_6$
$y_4$	$y_5$	$y_6$	$y_7$
$y_5$	$y_6$	$y_7$	$y_8$
$y_6$	$y_7$	$y_8$	$y_9$
$y_7$	$y_8$	$y_9$	$y_{10}$
$y_8$	$y_9$	$y_{10}$	?

$$y_8, y_9, y_{10} \rightarrow \hat{y}_{11}$$

$$y_9, y_{10}, \hat{y}_{11} \rightarrow \hat{y}_{12}$$

$$y_{10}, \hat{y}_{11}, \hat{y}_{12} \rightarrow \hat{y}_{13}$$

# Multi-step forecasting - direct strategy

---

$$y_1, y_2, y_3, y_4, y_5, y_6, y_7, y_8, y_9, y_{10} \rightarrow ?, ?, ?$$

X			y
$y_{t-2}$	$y_{t-1}$	$y_t$	$y_{t+1}$
$y_1$	$y_2$	$y_3$	$y_4$
$y_2$	$y_3$	$y_4$	$y_5$
$y_3$	$y_4$	$y_5$	$y_6$
$y_4$	$y_5$	$y_6$	$y_7$
$y_5$	$y_6$	$y_7$	$y_8$
$y_6$	$y_7$	$y_8$	$y_9$
$y_7$	$y_8$	$y_9$	$y_{10}$
$y_8$	$y_9$	$y_{10}$	?

X			y
$y_{t-2}$	$y_{t-1}$	$y_t$	$y_{t+2}$
$y_1$	$y_2$	$y_3$	$y_5$
$y_2$	$y_3$	$y_4$	$y_6$
$y_3$	$y_4$	$y_5$	$y_7$
$y_4$	$y_5$	$y_6$	$y_8$
$y_5$	$y_6$	$y_7$	$y_9$
$y_6$	$y_7$	$y_8$	$y_{10}$
$y_8$	$y_9$	$y_{10}$	?

X			y
$y_{t-2}$	$y_{t-1}$	$y_t$	$y_{t+3}$
$y_1$	$y_2$	$y_3$	$y_6$
$y_2$	$y_3$	$y_4$	$y_7$
$y_3$	$y_4$	$y_5$	$y_8$
$y_4$	$y_5$	$y_6$	$y_9$
$y_5$	$y_6$	$y_7$	$y_{10}$
$y_8$	$y_9$	$y_{10}$	?

$$y_8, y_9, y_{10} \rightarrow \hat{y}_{11}$$

$$y_8, y_9, y_{10} \rightarrow \hat{y}_{12}$$

$$y_8, y_9, y_{10} \rightarrow \hat{y}_{13}$$

# Multi-step forecasting - multi-output strategy

---

$$y_1, y_2, y_3, y_4, y_5, y_6, y_7, y_8, y_9, y_{10} \rightarrow ?, ?, ?$$

X			y		
$y_{t-2}$	$y_{t-1}$	$y_t$	$y_{t+1}$	$y_{t+2}$	$y_{t+3}$
$y_1$	$y_2$	$y_3$	$y_4$	$y_5$	$y_6$
$y_2$	$y_3$	$y_4$	$y_5$	$y_6$	$y_7$
$y_3$	$y_4$	$y_5$	$y_6$	$y_7$	$y_8$
$y_4$	$y_5$	$y_6$	$y_7$	$y_8$	$y_9$
$y_5$	$y_6$	$y_7$	$y_8$	$y_9$	$y_{10}$
$y_8$	$y_9$	$y_{10}$	?		

$$y_8, y_9, y_{10} \rightarrow \hat{y}_{11}, \hat{y}_{12}, \hat{y}_{13}$$

→ The multi-output strategy requires a model that can deal with multiple outputs, e.g. neural networks.

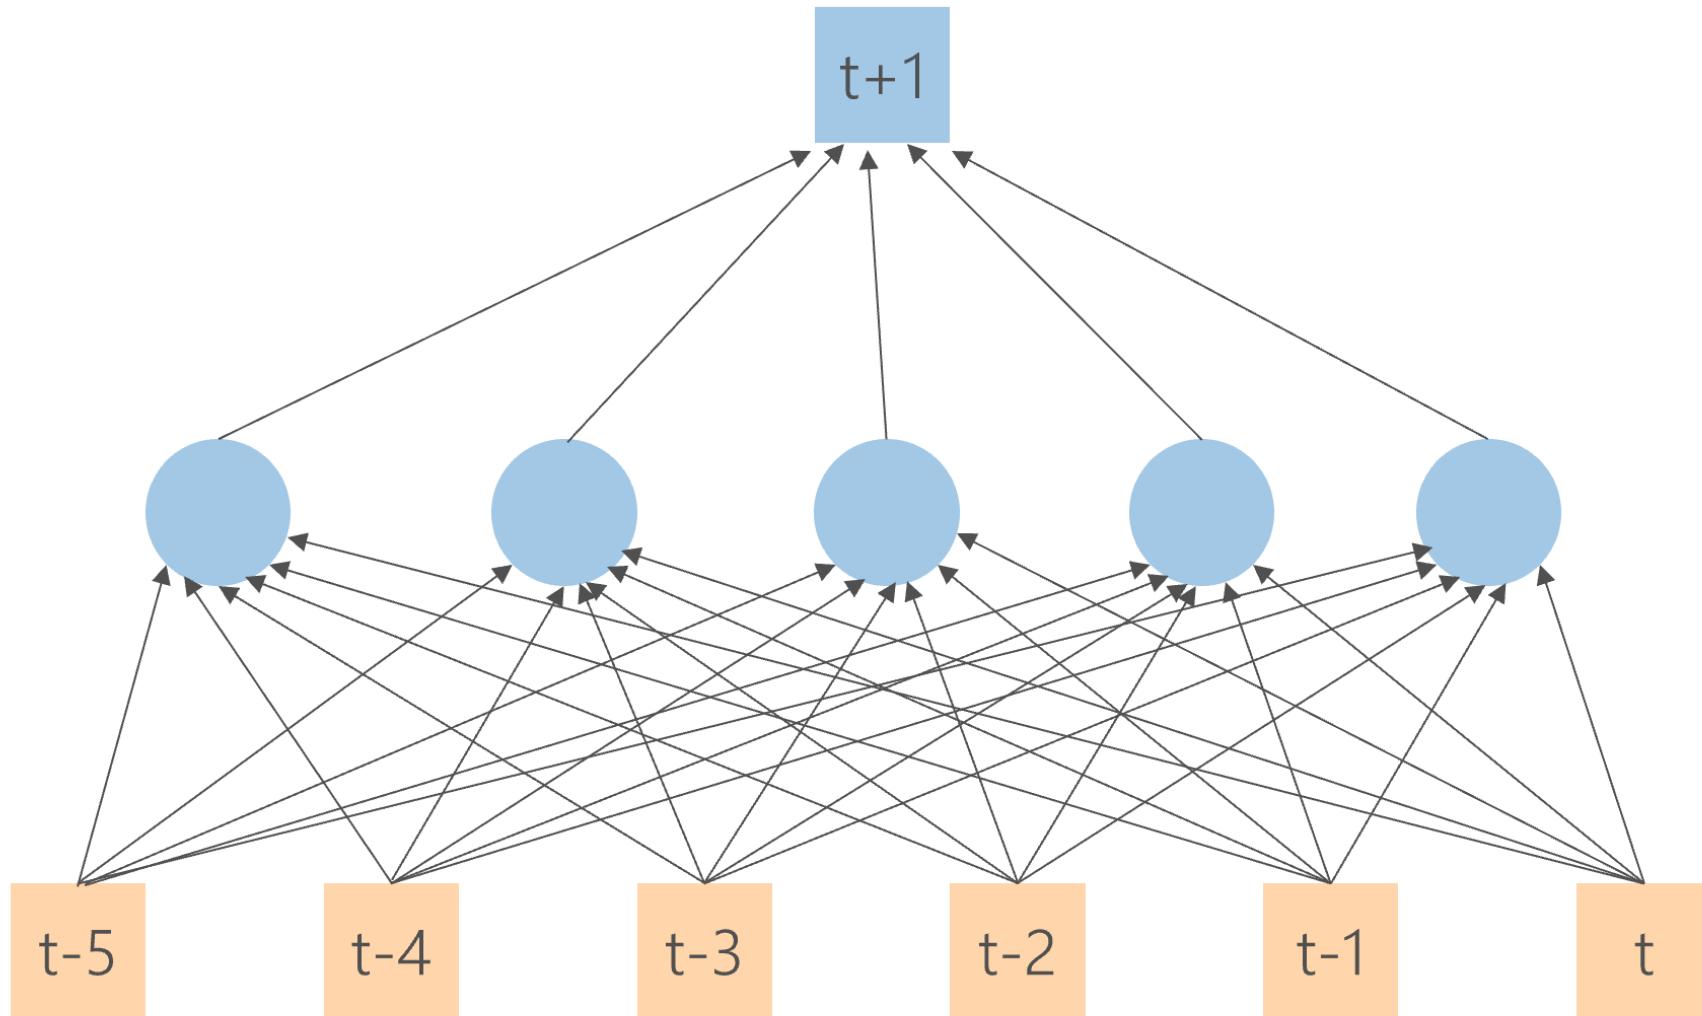
# Why (Deep) Neural Networks?

---

- Deep learning model has been shown to perform well in many scenarios
- Very effective at feature extraction
- Flexible and expressive
- Easily inject exogenous features into the model
- Learn from large time series datasets
- ...
- Challenges
  - Require a lot of data (in general)
  - Computationally demanding
  - Hard to train

# Single output network

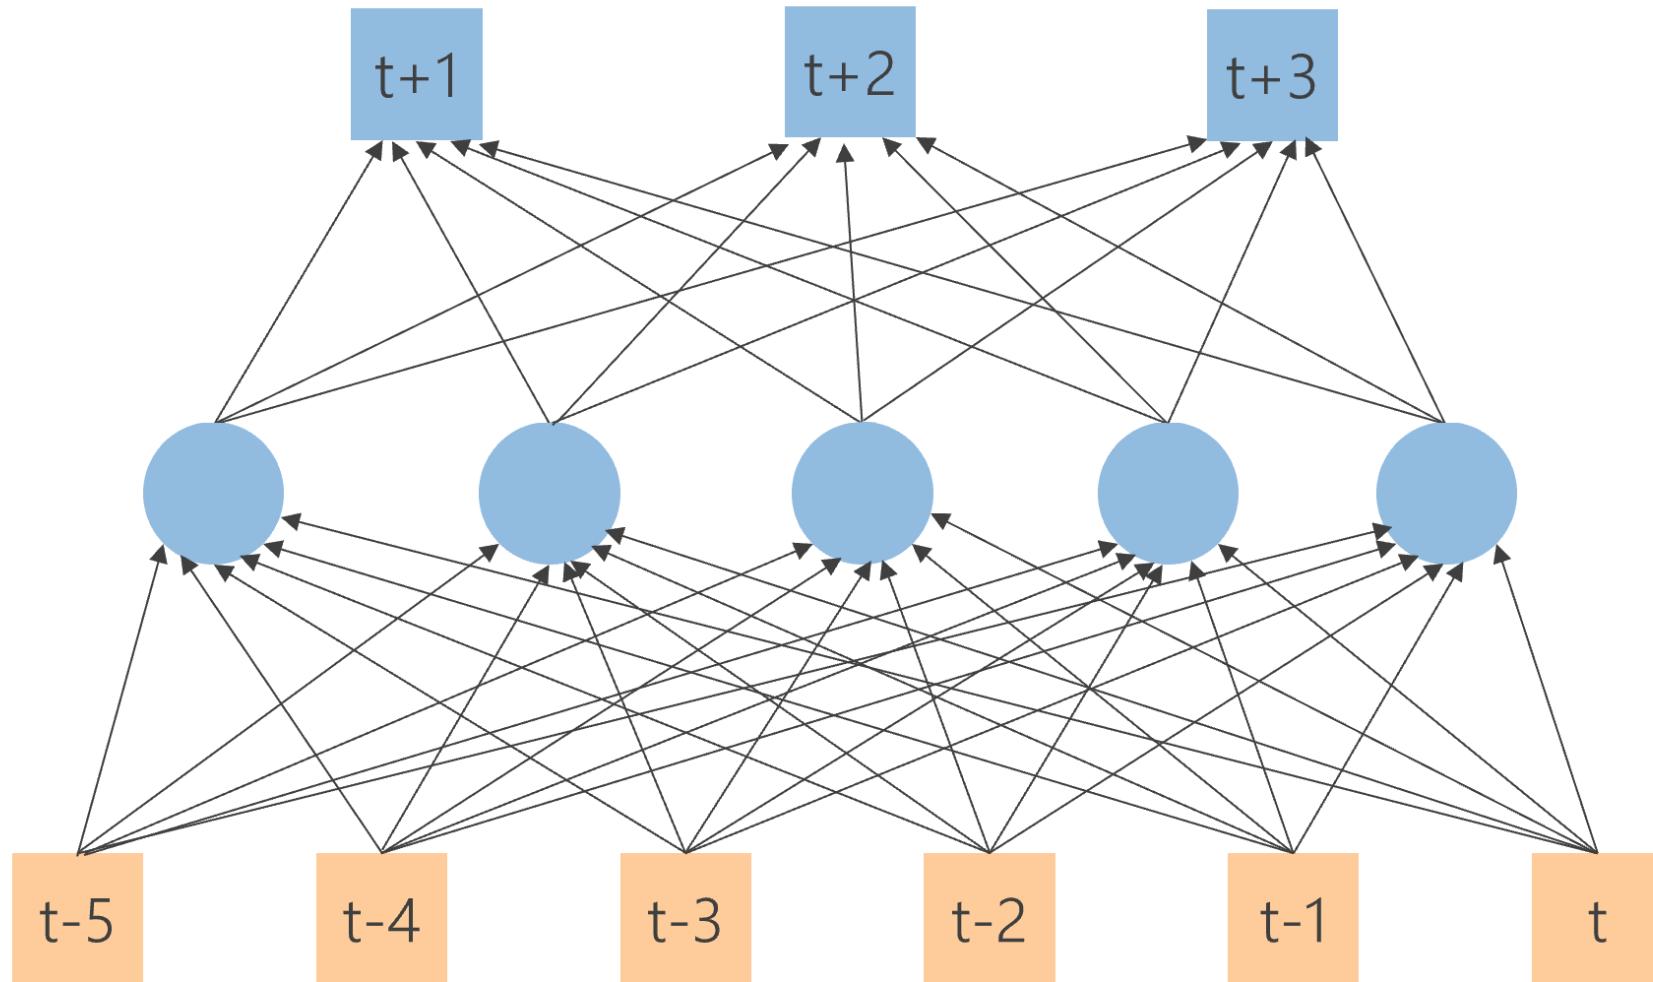
---



→ Multi-step forecasts are obtained recursively

# Multi-output network

---



→ The network produces multi-step forecasts.

# Neural network hyperparameters

---

Hyperparameters are adjustable parameters that define the model architecture and govern the learning process. By contrast, other parameters (such as node weights) are derived via model training.

- Architecture
  - Types of layers, number of layers, layer order, number of neurons per layer, layer activations, etc.
- Optimization
  - Optimizer, weight initialization, learning rate, batch size, number of epochs, stopping criterion, etc.
- Loss function
  - Loss function, form of regularization, etc.
- ...

# Hyperparameter search/tuning

---

- Search across various hyperparameter configurations
- Find the configuration that results in best (out-of-sample) performance
- Grid search vs random search
- Challenges
  - Huge hyperparameter space to explore
  - Computationally and time demanding